

# Sorting with Human Intelligence

Cole Kurashige  
Harvey Mudd College

March  
2020

## Abstract

Many comparison-based sorting algorithms have been introduced in years past, but none are capable of comparing elements of two different types. We present a novel algorithm called `Turksort` which uses human intelligence to sort lists with truly arbitrary contents. We also present an implementation that can be found at <https://github.com/cole-k/turksort>. We analyze its performance with respect to time, accuracy, as well as a novel metric called monetary complexity.

## 1 Background

Lower bounds for time and space complexity have been long-established for comparison-based sorting algorithms. Many sorting algorithms have been developed which vary in the trade-offs they make for these complexities. Little has been done, however, to examine what it means to make a comparison.

In statically-typed programming languages like Java or C++, comparisons such as equal to (`==`) or greater than (`>`) are often restricted to operating on two elements of the same type<sup>1</sup>. These languages consider it a compilation error to compare elements of differing types.

Dynamically-typed programming languages like Python do not consider it a compilation error; however, it is usually a runtime error to make these comparisons.

Listing 1: Comparisons in Python 3.7.6

```
>>> -1 > '-2'
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
TypeError: '>' not supported between instances of 'int' and 'str'
```

JavaScript, ever the staunch opponent of reason, will gladly compare two objects of differing types. But just because JavaScript can do something does not mean that JavaScript does it right.

---

<sup>1</sup>At least, without trickery or custom comparators.

It will correctly report that `-1` is greater than `"-2"`, but it erroneously considers `"-1"` less than `"-2"`. It also doesn't get that `"three"` is less than `"four"`, and it certainly does not know that `"one pound of feathers"` is just as heavy as `"one pound of bricks"`.

Listing 2: Comparisons in JavaScript (Node.js 12.12.0)

```
> -1 > "-2"
true
> "-1" > "-2"
false
> "three" < "four"
false
> "one pound of feathers" == "one pound of bricks"
false
```

Though it would be enjoyable to continue to mock JavaScript and its many questionable design choices, we cannot fault it much for these shortcomings.

JavaScript, like any programming language, interprets code. It treats queries like `"-1" > "-2"` as being a comparison on characters, not numbers, even if we humans can plainly see that JavaScript is being asked to compare negative one and negative two. But we cannot call JavaScript an idiot without calling it a savant. It can perform remarkably complex calculations in the blink of an eye or bring the fastest of hardware to a slow grind.

Computers are limited at processing much of the information that is so easy for us humans to immediately understand, like pictures of dogs or whether `"-1"` is greater than `"-2"`. And we are limited at processing much of the information that is so easy for computers to understand, like the exact colors of the millions of pixels in a picture of a dog or the product of two very large numbers.

Listing 3: Complex Operations in JavaScript (Node.js 12.12.0)

```
> 0.1 + 0.2
0.30000000000000004
> 1000000000 * 500000000
500000000000000000
```

Together, computers and humans can cover for each other's inadequacies, which is the premise of Human Intelligence Sorting. This type of sorting has yet to been realized in traditional sorting algorithms. At least until now.

## 2 Turksort

The idea behind Turksort is simple: let computers handle all of the sorting tedium and let humans handle all of the comparisons. It ends up being not that different from most sorting algorithms.

The algorithm differs only when two values need to be compared. When this happens, a form asking which of the two is greater<sup>2</sup> is generated. This

---

<sup>2</sup>"Neither" is an option, too.




Figure 1: How Turksort comparisons work.

form is sent to Amazon’s Mechanical Turk (MTurk)<sup>3</sup> where a worker (known as a “Turker”) fills it out. The answers is sent back and then used for the comparison. Figure 1 depicts this process.

An implementation of Turksort is available at <https://github.com/cole-k/turksort>. Any analyses in this paper will be with reference to this implementation. It is worth noting that Turksort refers to any algorithm that sorts using Turker-based comparisons, so other variants may be developed.

The implementation is a modification of quicksort. Quicksort works by selecting an element in the list (called the “pivot”) and comparing all of the other elements to it. It then partitions the list into three groups: those less than, equal to, or greater than the pivot. It recursively sorts all three partitions and combines them in order, producing a sorted array.

The partitioning process requires queries to be made comparing the pivot to each of the other elements in the list. These comparisons are collected and sent to MTurk for evaluation. This allows us to batch the computations, since querying MTurk is slow in comparison to regular comparison. The answers from MTurk are then used in the partitioning, and the algorithm proceeds as usual.

### 3 Analysis

In this section, we analyze the performance of Turksort (section 3.1) as well as its accuracy (section 3.2 on the following page).

#### 3.1 Performance

Turksort is not an algorithm whose performance should be measured by traditional means. It, however, can be.

Since it retries until it gets a response from the Turker, Turksort technically has an unbounded time complexity. Even assuming that the response time from the Turker is bounded and proportionate to the query size, the asymptotic time complexity is that of quicksort. The average response time, even for shorter

<sup>3</sup><https://www.mturk.com/>

queries, is around 5 minutes, so the constants on the time complexity are very large.

The novel performance measurement we propose for Turksort is a cost metric. We call it monetary complexity. This is the asymptotic cost of performing a computation. Because one currency differs from other currencies by a scaling factor, the monetary complexity’s monetary base does not matter, much like logarithmic base in asymptotic complexity does not matter. We use a monetary base of USD.

It takes a Turker about 1 second to answer a single query. Since minimum wage in California is presently \$12.00, we pay 1 cent for every 3 queries a Turker answers, paying a floor of 1 cent if they are answering fewer than 3. We measured the monetary complexity of Turksort with respect to the number of elements in the list, with lists up to size 10000. We calculated the cost as being the average of five trials. Because the authors do not have any grant money, testing was done using a simulated Turker.

We introduce a new notation  $\$(f(n))$  to denote monetary complexity: it means that a computation has cost asymptotically proportionate to  $f(n)$ . Turksort has a monetary complexity of  $\$(n \log n)$ ; other common sorting algorithms have an effective monetary complexity of  $\$(1)$ <sup>4</sup>. You can observe this complexity in figure 2 on the next page.

It is evident that Turksort should only be used in cases where traditional computing does not have sufficient intelligence. The tradeoffs for using Turksort are both time and money, although common adages suggest that this is only a single tradeoff. In section 4 on the following page we discuss potential solutions to these tradeoffs. As it turns out, there is a third, unexpected tradeoff, which is accuracy. We discuss this below.

## 3.2 Accuracy

Surprisingly, Turksort is not a deterministic algorithm. This is because humans are not deterministic<sup>5</sup>. Not only is Turksort nondeterministic, it is also sometimes wrong. This is because Turkers do not always perform the right computations. Even on simple queries, such as  $2 > 3$ , they can give an incorrect answer.

This does not mean that Turksort is a useless sorting algorithm. There is a simple tradeoff between accuracy and speed: the less time a Turker spends answering a question, the more likely it is to be incorrect. Turksort is already not winning any races, and that is fine since it serves a specific purpose that regular sorting algorithms do not. So making it slightly slower for greater accuracy is a worthwhile tradeoff. We discuss how to mitigate the problem of accuracy by making more or slower queries in section 4 on the next page.

---

<sup>4</sup>Though they cost money by way of using electricity, this is a negligible cost and can be considered effectively constant.

<sup>5</sup>Although it is unknown whether individual humans are deterministic, in general no two humans perform comparisons exactly alike.




Figure 2: The monetary complexity of Turksort plotted for lists of size up to 10000.

## 4 Future Work

In the previous section, we discussed some limitations of Turksort. In this section we will discuss how these limitations may be overcome. Section 4.1 discusses ways to improve its accuracy and section 4.2 on the next page discusses ways to improve its performance. Turksort is very widely applicable and useful, so we do not need to mention potential applications or uses.

### 4.1 Improving Accuracy

The most important problem Turksort presently faces is an accuracy issue. There are two potential solutions.

First, Turkers could be forcibly slowed down by imposing a time limit before they can answer a comparison. This will prevent them from answering so fast that they get it wrong. The bulk of the time spent waiting in Turksort is in waiting for a Turker to start responding, so this will not extend the duration of the algorithm significantly, especially for shorter queries.

Second, Turksort could issue multiple requests for the same query. This way, majority voting from the Turkers could be used to increase the accuracy. Because these queries would be sent out in parallel, it is unlikely that this will have significant impacts on time. This has the additional benefit of making it easier for Turksort to sort so-called “trick comparisons,” like a query of “**a pound of bricks**” > “**a pound of feathers**”. If, during a computation, a query is suspected of being a “trick comparison,” the algorithm can take the minority response instead.

## 4.2 Improving Performance

Performance is less important for Turksort, given the time it takes to answer queries, but as the field of Human Intelligence Sorting grows, faster and cheaper variants will become more useful.

One way of improving performance is parallelization. Since queries take a long time to answer, Turksort could issue multiple queries at once. This can either be realized by modifying the underlying sorting algorithm to be more parallel, by making all  $\binom{n}{2}$  comparison queries at once (thereby increasing the  $\$(n \log n)$  monetary complexity to  $\$(n^2)$ ), or by performing “branch prediction” and guessing what the next queries might be.

An obvious way of reducing the monetary complexity is to slow the rate at which Turkers are rewarded. Though it might seem illegal to not pay Turkers minimum wage, Turksort does not need to pay its Turkers since the gratification that they are advancing human progress is payment enough. However, without large constants, we believe  $\$(1)$  Turksort algorithms to be impossible, as Turkers are not motivated by this gratification. We are presently exploring a  $\$(\log^2 n)$  variant of Turksort.

## 5 Acknowledgements

We would like to acknowledge Arya Massarat, Andrew Pham, and Giselle Serate for testing the algorithm for  $\$(1)$  monetary cost. We would also like to acknowledge all of the Turkers who tested the paid version of the algorithm. And we would finally like to acknowledge Andrew Pham, Giselle Serate, and Max Tepermeister for proof reading this paper.

Finally, we would like to acknowledge a blog post by Mikey Levine describing a similar idea with the same name<sup>6</sup> for teaching us to search the internet more carefully after we come up with so-called “novel” ideas and then write papers on them. Indeed, careful inspection reveals that this general idea has been explored a few times prior, although thankfully not in the same ways as this paper.

---

<sup>6</sup>[http://games.hazzens.com/blog/2014/02/27/turk\\_sort.html](http://games.hazzens.com/blog/2014/02/27/turk_sort.html)