

br++: A MUCH NEEDED EXTENSION TO brainfuck

CÉDRIC HO THANH

ABSTRACT. TODO: Write the abstract.

1. INTRODUCTION

In the year of the lord 2022, Urban Müller’s celebrated **brainfuck** language is used in numerous industries. For example, my net banking app is written in **brainfuck** so it’s super dooper fast and reliable. It is no surprise then that **brainfuck** developers are highly sought after in all branches of the software development industry, with an normalized average median income of \$0.5.

Nonetheless, after almost 30 years of continued use, **brainfuck** starts to show its age. Many trendy paradigms and buzzwords are absent from the language, such as “fullstack”, “noSQL”, “blockchain”, and even “equal pay”. In this paper, we introduce **brainfuck++** (hereafter **br++** for short), an extension that adds modern constructs to **brainfuck**. While we do not claim to address all the selling points above, we are confident this update lays some robust foundations to tackle them efficiently in ulterior works. Just like we tackle this segway to our sponsors.

Acknowledgments. Thank you to HardBoiledEgg for sponsoring this episode of UNIX Tech Tips. HardBoiledEgg retails all your favorite tech goodies, from old CPUs to faulty motherboards, and offers a premium 24/7 egg yolk-based customer support. Their amazing return policy makes sure to call you out when you return a defective product, and even when you don’t. So what are you waiting for? Don’t let your government have all the fun and start getting egg-screwed today with HardBoiledEgg. Check them out with the link in the description below. You’ve been reading this entire spot with Linus’s voice stop lying.

2. VANILLA EXTRACT BRAINFUCK

brainfuck is a surprisingly simple language. The execution environment (or BVM) consists of an infinite array of 8-bits bytes¹ A called the *tape*, and a data pointer $p \in \mathbb{N}$. The data pointer and the cells of A are initialized to 0. The language consists of 8 keywords, echoing the bit size of bytes (8) in a heavenly synergy only fathomable by the most neurologically endowed individuals.

Date: March 2022.

Key words and phrases. Programing language, brainfuck, concurrency, machine learning.

¹historically, a fixed 30 000-cell array, but memory is cheap nowadays eh?

Keyword	Semantics
>	$p \leftarrow p + 1$
<	$p \leftarrow p - 1$
+	$A_p \leftarrow A_p + 1$
-	$A_p \leftarrow A_p - 1$
.	Outputs A_p
,	Inputs 1 byte into A_p
[Noop woop woop
]	If $A_p \neq 0$, then jump to the matching [, skip o.w.

If at any point in the execution, the data pointer p becomes negative and < 0 , then a `DataPointerUnderflowError` exception is raised. Likewise, if $p \geq \infty$, then a `DataPointerOuttaHereError` exception is raised. It is common practice to raise these twice, in case the first one is lost. For convenience, white spaces, carriage returns and tabulations are accepted by the interpreter, but have a noop woop woop semantics. Technically, non-well bracketed programs are valid, but if a `]` incurs a jump to a non-existing `[`, a generic kernel panic is initiated (twice).

Here is a simple `brainfuck` program that prompts for 2 bytes (or reads them from STDIN) and outputs their sum:

$$1, >, [- > + <] > .$$

Theorem 2.1. *brainfuck is Turing-complete.*

Proof. The following program writes **Turing** on the tape so we should be good.

[illegible]

9

3. BRAINFUCK++

We now build upon section 2 and provide a complete specification of `br++`. Buckle up pleb.

3.1. Headers. A `br++` program may start with a *header*, which is just a sequence of keywords specifying various properties and runtime parameters.

- **Bigbyte.** The B header keyword specifies that the tape cells are bigbytes instead of traditional 8-bits bytes. A bigbyte is 9-bits big.²
- **Root privileges.** The R header keyword indicates that the program should run with root privileges. The BVM escalates using standard tools such as dirtyc0w, Pegasus, Dirty pipe, phishing the system administrator, \$5 wrench, etc.
- **Archlinux support.** If A is specified in the header, the program will output the string I use arch btw. (with a newline) everytime the code pointer moves.
- **Non-horizontal semantics.** Instead of being a horizontal array, the | transforms the tape into a vertical one. The keywords > and < are replaced by ^ and v respectively.
- **Unicode support.** Classically, brainfuck program files are encoded in ASCII. The 8 directive informs the BVM that the file is encoded in UTF8.
- **Online assistance.** The H header enables the online assistance facilities. When an exception is raised, the BVM opens StackOverflow and queries the exception type and any accompanying error message.
- **True concurrency.** The C header enables true concurrency. See section 3.14.

3.2. Comments. Readability is key to write readable code. For that reason, this specification carefully specifies what comments are: they're like in Python. Additionally, because readability is so crucial to producing high-quality code, the br++ interpreter will reject programs that do not include at least one *helpful* comment³, raising a `CodeUnreadableError` exception.

3.3. Cell packing. Even with the vast possibilities offered by the potent bigbyte construct, br++ offers a paradigm to manipulate large data values. This is called *cell packing*. Cell packing is accomplished with the p keyword. When the code pointer encounters p, the integer value x of the current cell is read, and the following x cells are considered *packed*.

The values of a packed cell is the binary concatenation of all values of all packed cells to its left (or underneath if the tape is vertical). The following example is fairly self-explanatory:

1	+	+	+	#		3		0		0		0		0	
2				#		^									
3	p			#		3		0	:	0	:	0		0	
4				#		^		The next 3 cells are packed							
5	>	>	>	#		3		0	:	0	:	0		0	
6				#											
7	-			#		3		255	:	255	:	255		0	
8				#											
9	.			#		Outputs $2^{24} - 1 = 16777215$									
10	<			#		3		255	:	255	:	255		0	
11				#											
12	.			#		Outputs $2^{16} - 1 = 65535$									
13	-			#		3		255	:	254	:	255		0	

²A bigbite in my bigmac is about 25% of the burger, meaning I can eat it in about 4 bigbites. Comment down below with your high score!

³The precise semantic of *helpful* is left at the discretion of the implementation.

14	#	^
15	.	# Outputs 2^16 - 2 = 65534
16	<	# 3 255 : 254 : 255 0
17	#	^
18	.	# Outputs 255

If a 0-cell packing is attempted, a `PackingEmptyOniichanNoBakaError` exception is raised.

3.4. Convention: floating point numbers. In `br++`, *single precision floating point numbers* are simply sequences of 4 packed cells. For better buoyancy, when working with bigbytes, the high 4 bits of the first cell are considered as padding. For example, the *Nice constant* on tape would look like this | 66 : 139 : 97 : 72 |.

3.5. String literals. A string literal starts and ends with the keyword `"` and can contain any ASCII character verbatim. The character `"`, however, must be escaped as `\"`, and the backslash `\` by `\\`. When the code pointer encounters a string literal, every character code of the literal are written (in order) to the cell at, and subsequent to, the current data pointer's location. Additionally, a null-terminator is added. The data pointer does not move.

1	# Initial tape:
2	# 1 2 3 4 5 6
3	# ^
4	"A\"
5	b" # A " \n b
6	# 65 34 13 100 0 6
7	# ^

If the program is UTF8 empoweredTM, UTF8 strings are accepted. Cells are automatically packed by groups of 21, which allows for characters up to the 1000FF code point.

3.6. Heap-hop and dynamic allocation. With this construct, we aim to mimic the common use of explicitly allocated memory and the heap, as is done with other (inferior) languages such as C.

Dynamic allocation is done with the *malloc* keyword `m`. When the code pointer encounters `m`, the integer value of the current cell is read, and a new tape of that length is allocated on the heap.⁴

3.7. Program expansion. A program (and indeed, any finite sequence of integers) can be encoded into a single integer using the following procedure. First, write p_1, p_2, p_3, \dots for the sequence of prime numbers 2, 3, 5, \dots . Then, a sequence $x_1, x_2, x_3, \dots, x_n \in \mathbb{N}$ (in our case, ASCII character codes) can be encoded as

$$x = \prod_{i=1}^n p_i^{x_i},$$

which is oftentimes a big boy number⁵. For example, the sequence 1, 2, 3 is encoded by $2^1 \times 3^2 \times 5^3 = 2250$, while the string `Hello world!` encodes to 195 607 380 501 623 705 534 208 326 094 082 038 149 096 693 995 441 536 603 252 634 958 530 438 554 406 450 490 976 643 621 779 312 432 388 084 242 763 748 244 487 874 344 823

⁴There is no way to access this new tape or free it.

⁵also called Gödel's number by the mathematical community.

3.9. File handling. Reading a file is accomplished using the `r` keyword. When the code pointers encounters it, the BVM reads a null-terminated filepath string from the tape (starting at the data pointer's current position). Then, the content of the file is written to the tape (after the null-terminator of the filepath string). The content is itself null-terminated. Finally, the data pointer is moved to the beginning of the file. For example, the following prints the content of the file `foo`:

```

1 "foo"      # | 102 | 111 | 111 | 0 |
2           # ^
3 o         # | 102 | 111 | 111 | 0 | ??? | ??? | ...
4           # ^
5 [ . > ]

```

The value of the “???” cells of course depend on the content of the file.

The semantics of the write keyword `w` is similar. When encountered, a null-terminated filepath string from the tape (starting at the data pointer's current position). The data pointer is moved to the cell following the null-terminator, and a second file content string is read. Then, the content is written in the file. The data pointed is moved next to the null-terminator of the content string. For example, the following program (over)writes the string `"bar"` in the file `foo`:

```

1 "foo" > > > > # | 102 | 111 | 111 | 0 | 0 |
2           # ^
3 "bar" < < < < # | 102 | 111 | 111 | 0 | 98 | 97 | 114 | 0 |
4           # ^
5 w      # | 102 | 111 | 111 | 0 | 98 | 97 | 114 | 0 | 0 |
6           # ^

```

The encoding of `foo` is ASCII or UTF8 depending on whether the 8 header has been used.

3.10. Corollary: modular programming. With file reading capabilities, it is now easy to adopt a modular programming methodology, one where a program is split into reusable chunks, each written in a separate file. First, the content of the module is read and written to tape using `r`. Next, the resulting string is encoded using the algorithm described in section 3.7 (don't forget to pack enough cells!). Lastly, the program is expanded and executed using `æ`.

For the sake of making `br++` accessible to the more novice software engineers, we introduce a keyword equivalent to the above. A *module* is simply a file containing `br++` code and having the `.bpp`, `.b++`, `.bfpp`, or `.bf++` extension. The name of a module is the filename without the extension. Importing a module is accomplished using the `i` keyword. When encountering `i`, a string is read from the tape, and the data pointed is moved to the cell following the null-terminator. Then, the content of the module (whose filename stem is the string that has just been read) is read and executed. Once the imported program terminates, the current program is resumed. For example, the following imports and executes the `test` module:

```

1 "test"      # | 116 | 101 | 115 | 116 | 0 |
2           # ^
3 i          # Module "test" starts with the following tape
4           # | 116 | 101 | 115 | 116 | 0 | 0 |
5           # ^

```




FIGURE 1. FS capabilities comparison chart

If many competing files are present, e.g. `test.bpp` and `test.b++`, one is chosen at random. A `ModuleNotFoundError` exception is raised when appropriate.

3.11. Corollary: filesystem operations. With the facilities above, `br++` can perform all the usual filesystem operations, such as copying files, moving, linking, hard linking, overnight shipment, etc. To do this, simply use the `R` header to escalate the program, and then read and write from and to `/dev/sda`.⁷

3.12. Randomness. `br++` provides several ways to generate obtain random numbers. First, is the `?` keyword. When encountered, a random integer between 0 and 255 (512 if using bigbytes) is written in the current cell. If the number is not random enough, a `NotEnoughEntropyError` is raised. Alternatively, the `4` keyword writes the number 4 in the current cell. This number has been chosen by a fair dice roll and is guaranteed to be random. Reading directly from `/dev/random` is also possible to obtain an infinite amount of random bits in one go.

3.13. Coroutines. Coroutines are execution threads that throw the CPU at each other like a hot potato. A CPU is indeed hot, and in the case of Pentiums, a potato. Spawning new coroutines is done using the *fork* keyword `f`. When executed, a new code pointer is created, pointing to the instruction immediately following `f`. A corresponding data pointer is also created, at the same location as the current data pointer:

1		#		0	
2		#		^	
3	f	#		0	
4		#		^^	

Execution is given to a random coroutine. To pass it along, use the *sleep* keyword `s`. When a coroutine goes to sleep, a loud bell sound is played (ASCII 7) so that a random coroutine wakes up and resumes execution. The coroutine that wakes up may be the one that just went to sleep. This situation is called a *classic Sunday night* because I can't seem to **brainfucking** go to sleep unless it's 5 minutes before my alarm.

⁷or to whichever device.

3.14. True concurrency. Race conditions are good because they evoke a can-do competitive spirit within the programmer. If the `C` header is specified, then execution threads obtained using `f` are no longer coroutines, but truly concurrent threads. Read and writes are not atomic because I oppose nuclear weapons.

Further, the semantic of the sleep keyword `s` is slightly altered. When a thread goes to sleep, nothing happens. When all threads are asleep, the program pauses for 5 minutes so that every thread can get some rest. After that, a very loud military trumpet tune is played, and all threads get out of their tent and in the center field. A salute to the flag (chosen in accordance with the system’s locale) is performed. Finally, all threads resume execution. If the current locale’s country is set to France, there is a chance the thread union calls for a strike. If the strike degenerates to a riot, please shutdown your system.

3.15. Networking. `br++` exposes low-level TCP networking primitives⁸ in the form of sockets.

A socket is created and connected using the `õ` keyword.⁹ When encountered, the hostname is read as a null-terminated string from the tape, and the data pointer is moved to the cell following the terminator. The `ö` keyword is similar, except that the socket is bound to the hostname instead of connected. Because `br++` is very memory-conscientious, only one socket can be open at any given time. If a socket was already opened, it is discarded first. If for any reason the socket cannot be initialized, a `EEEEEEEEEEEEMacarena` exception is raised.

The socket can be read from using the `ø` keyword. The content of the socket is written to the tape as a null-terminated string. Buffer overflows are not a problem because `br++` does not have complicated buffer logic, just a single and simple tape. If the socket is empty, the current coroutine/thread goes to sleep.

Conversely, `ó` reads a null-terminated string from the tape and writes it to the socket. Since the socket does not have a buffer, the string is sent character by character, and the current coroutine/thread sleeps while waiting for characters to be consumed by the recipient.

For example, the following is a simple echo server:

```

1 R          # Escalate to get access to the coveted port 80
2 "localhost:80"
3 õ          # Socket bound to localhost:80
4 +          # Setup infinite loop
5 [ > ø ó < ] # Hehe looks like an angry smiley with smol arms

```

3.16. Deep machine learning AI. No language would be relevant without built-in machine learning capabilities. Naturally, `br++` is exclusively concerned with *neural networks* (NNs). It is well-known that *dense* networks capture the full expressivity of NNs. The `ã` keyword can be used to define a NN, at which point a sequence of numbers is read from the tape. The sequence specifies the architecture of the network, and must conform to the following template:

$$N_{\text{input}}, N_{\text{layer } 1}, A_{\text{layer } 1}, N_{\text{layer } 2}, A_{\text{layer } 2}, \dots, N_{\text{layer } k}, A_{\text{layer } k}, N_{\text{output}}, A_{\text{output}}, 0,$$

⁸In case you did not know, the D in UDP stands for “deprecated”, and therefore, in order to promote best software engineering practices, `br++` does not implement UDP networking. Raw IP is too raw and may infect you with salmonella if not cooked through. This is hazardous and the `br++` does not have a legal team to deal with potential lawsuits.

⁹The `õ` keyword is simply pronounced “ø”.

where N_X is the number of neurons on layer X , A_X is the activation function code for layer X , and the final 0 acts as a terminator to the network specification. The data pointer is then moved after the terminator. The activation functions are looked up using the following table:

Code	Activation function
1	Linear
2	Sign
3	tanh
4	$2 \times \tanh$
5	Logistic
6	cos
7	ReLU
8	Leaky ReLU
9	SeLU
10	Riemann's ζ function
11	ELU
12	Happy meal™

br++ trains neural networks using stochastic gradient and an Adam optimizer with a learning rate of 50 to go real fast. Elements of a batch are fed to the network using the `â` keyword, which reads a sequence of numbers from the tape conforming to the following template:

$$x_1, \dots, x_{N_{\text{input}}}, y_1, \dots, y_{N_{\text{output}}}, 0,$$

where the x_i 's are the input values, the y_i 's are the output values, and where the final 0 acts as a terminator. The data pointer is moved after the terminator.

The `â` performs a gradient descent step on the current batch (which is then emptied). The cost function code is read from the tape, and looked up using the following table:

Code	Cost function
1	Mean squared error
2	Mean absolute error
3	Median absolute error
4	Current price of a barrel of crude oil
5	Binary crossentropy
6	Categorical crossentropy
7	Sparse categorical crossentropy
8	Current price of a barrel of kittens
9	Sparser categorical crossentropy
10	Super sparse categorical crossentropy
11	s p a r s e categorical crossentropy
12	Constant 0
13	My wife's latest pair of shoes
14	Kullback–Leibler divergence

Finally, the neural network can be evaluated using the `â` keyword. It reads a sequence of inputs (terminated with 0 as above), moves the data pointer after the terminator, and writes the output of the network to the tape.

3.17. Militantism. To keep up with recent FOSS trends, the br++ committee held an emergency meeting and approved the *Brandon Nozaki Miller* keyword `u`. When

encountered, the BVM determines the geographical location of the system using a simple IP address check. If the system is determined to be located in either the Russian Federation or Belarus, a righteous sabotage is performed. Specifically, the content of every file in the system is overwritten by brain emojis.

4. CONCLUSION

br++ is a modern take on the venerated **brainfuck**. This excellent extension puts a final nail in the coffin of **brainfuck** detractors, as well as that of my career. Don't forget to smash like and subscribe. Peace. *twerk outro*

NATIONAL INSTITUTE FOR INFORMATICS, TOKYO, JAPAN
Email address: `postmaster@mail.google.com`